Advanced Database Systems, CSCI-GA.2434-001
New York University, Fall 2018

instructor: Dennis Shasha

shasha@cs.nyu.edu

212-998-3086

Courant Institute

New York University

251 Mercer Street

NY, NY 10012 USA

Office Hours: On Tuesdays before class in the Forbes building

November 12, 2018

1   Goals

To study the internals of database systems as an introduction to research and as a basis for rational performance tuning.

The study of internals will concern topics at the intersection of database system, operating system, and distributed computing research and development. Specific to databases is the support of the notion of transaction: a multi-step atomic unit of work that must appear to execute in isolation and in an all-or-nothing manner. The theory and practice of transaction processing is the problem of making this happen efficiently and reliably.

Tuning is the activity of making your database system run faster. The capable tuner must understand the internals and externals of a database system well enough to understand what could be affecting the performance of a database application. We will see that interactions between different levels of the system, e.g., index design and concurrency control, are extremely important, so will require a new optic on database management design as well as introduce new research issues. Our discussion of tuning

will range from the hardware to conceptual design, touching on operating systems, transactional subcomponents, index selection, query reformulation, normalization decisions, and the comparative advantage of redundant data. This portion of the course will be heavily sprinkled with case studies from database tuning in biotech, telecommunications, and finance.

Because of my recent research interests, this year will include frequent discussions of "array databases," the extension of relational systems to support ordered data such as time series in finance, network management etc. You will even get to program in our system aquery: https://github.com/josepablocam/aquery2q

2 Mechanics

YOU MUST BE ENROLLED IN THIS CLASS TO SIT IN ON THE LECTURES.

2.1 Texts and Notes

The first text will be used for the first half of the course and the second text in the second half. The notes will be used throughout the course.

- *Concurrency Control and Recovery in Database Systems* by Bernstein, Hadzilacos, and Goodman, Addison-Wesley, 1987. ISBN 0-201-10715-5 Available for free from Phil Bernstein's website at Microsoft. You can find a copy in this current directory at ccontrol.zip

There are also four optional books which are very nicely written:

- *Database Tuning : Principles Experiments and Troubleshooting Techniques* Dennis Shasha and Philippe Bonnet, Morgan Kaufmann Publishers, June 2002, ISBN 1-55860-753-6, Paper, 464 Pages.

- Transaction Processing: Concepts and Techniques by Jim Gray, Andreas Reuter 1002 pages ; Publisher: Morgan Kaufmann; 1st edition (1993) ISBN: 1558601902

- *Transactional Information Systems: Theory, Algorithms, and the Practice of Concurrency Control and Recovery* Gerhard Weikum, Gottfried

Vossen The Morgan Kaufmann Series in Data Management Systems, Jim Gray, Series Editor May 2001, 944 pages $79.95, ISBN 1-55860-508-8 If you want to be a world authority on these topics.

2.2   Prerequisites

Fundamental Algorithms I plus Data Base Systems I or equivalent (first 6 chapters of Ullman). If you don't have the database prerequisites, then you may take the course, but you must be responsible for understanding material covered in Database I: a reading knowledge of SQL and basic familiarity with indexes and third normal form.

2.3   Course Requirements

problem sets (40%), project (60%).

LATE HOMEWORKS OR PROJECTS WILL NOT BE ACCEPTED without a note from your physician or from your employer. (We will discuss the solutions on the day you hand in the assignment. That's why I don't want any late homeworks. As for projects, this is a question of fairness.)

On the other hand, collaboration on the problem sets IS allowed. You may work together with one other student and sign both of your names to a single submitted homework. Both of you will receive the grade that the homework merits. So, you may work alone or in a team of two, but no team larger than two. Please choose your partner carefully.

3   Syllabus — times are estimated

1. Overview of transaction processing, distributed systems, and tuning (1 week)

2. Principles of concurrency control for centralized, distributed, and replicated databases. (3 weeks)

3. Principles of logging, recovery, and commit protocols. (3 weeks)

4. Database Tuning (7 weeks)

```
        Tuning principles.
        Hardware, operating system,  and transaction subsystem
        Transaction Chopping
        Index tuning
        Tuning relational systems
        Tuning data warehouses
        Troubleshooting
        Case Studies from Wall Street and Elsewhere
```

5. Special topics: array databases, special indexes, time series.

4 Project

Your project is due on December 4, 2018. It will be graded by the last class at which point you will have nothing more to do. You have three possible projects:

1. Distributed replicated concurrency control and recovery. You may do this in a team of two.

2. An experimental study of tuning issues on a real system.

3. A paper on the topics I invite.

4.1 Possibility 1 — Replicated Concurrency Control and Recovery (RepCRec for short)

In groups of 1 or 2, you will implement a distributed database, complete with multiversion concurrency control, deadlock detection, replication, and failure recovery. If you do this project, you will have a deep insight into the design of a distributed system. Normally, this is a multi-year, multi-person effort, but it's doable because the database is tiny.

Data

The data consists of 20 distinct variables x1, ..., x20 (the numbers between 1 and 20 will be referred to as indexes below). There are 10 sites numbered 1 to 10. A copy is indicated by a dot. Thus, x6.2 is the copy of variable x6 at site 2. The odd indexed variables are at one site each (i.e. 1

+ (index number mod 10)). For example, x3 and x13 are both at site 4. Even indexed variables are at all sites. Each variable xi is initialized to the value $10i$ (10 times i). Each site has an independent lock table. If that site fails, the lock table is erased.

Algorithms to use

Please implement the available copies approach to replication using strict two phase locking (using read and write locks) at each site and validation at commit time. A transaction may read a variable and later write that same variable as well as others. Please use the version of the algorithm specified in my notes rather than in the textbook, just for consistency. Note that available copies allows writes and commits to just the available sites, so if site A is down, its last committed value of x may be different from site B which is up.

At each variable locks are acquired in a first-come first-serve fashion, e.g. if requests arrive in the order rlock1(x) rlock2(x) wrlock3(x) rlock4(x), then, assuming no transaction aborts, first T1 and T2 will share read locks on x, then T3 will obtain a write lock on x and then T4 a read lock on x. Note that rlock4 doesn't skip in front of wlock3, as that might result in starvation for writes.

Detect deadlocks using cycle detection and abort the youngest transaction in the cycle. This implies that your system must keep track of the oldest transaction time of any transaction holding a lock. (We will ensure that no two transactions will have the same age.) There is never a need to restart an aborted transaction. That is the job of the application (and is often done incorrectly). Deadlock detection need not happen at every tick. When it does, it should happen at the beginning of the tick.

Read-only transactions should use multiversion read consistency.

Test Specification

When we test your software, input instructions come from a file or the standard input, output goes to standard out. (That means your algorithms may not look ahead in the input.) Each line will have at most a single instruction from one transaction or a fail, recover, dump, end, etc. Some of these operations may be blocked due to conflicting locks. We will ensure that when a transaction is waiting, it will not receive another operation.

If an operation for T1 is waiting for a lock held by T2, then when T2 commits, the operation for T1 proceeds if there is no other lock request ahead

of it. Lock acquisition is first come first served but several transactions may hold read locks on the same item. So, if x is currently read-locked by T1 and there is no waiting list and T2 requests a read lock on x, then T2 can get it. However, if x is currently read-locked by T1 and T3 is waiting for a write lock on x and T2 subsequently requests a read lock on x, then T2 must wait for T3 either to be aborted or to complete its possesion of x; that is, T2 cannot skip T3.

The execution file has the following format:

begin(T1) says that T1 begins

beginRO(T3) says that T3 begins and is read-only

R(T1, x4) says transaction 1 wishes to read x4 (provided it can get the locks or provided it doesn't need the locks (if T1 is a read-only transaction)). It should read any up (i.e. alive) copy and return the current value (or the value when T1 started for read-only transaction). It should print that value in the format
x4: 5
on one line by itself.

W(T1, x6,v) says transaction 1 wishes to write all available copies of x6 (provided it can get the locks on available copies) with the value v. So, T1 can write to x6 only when T1 has locks on all sites that are up and that contain x6.

dump() gives the committed values of all copies of all variables at all sites, sorted per site with all values per site in ascending order by variable name, e.g.
site 1 – x4: 5, x2: 6, x3: 2, ... x20: 3
...
site 10 – x9: 5, x2: 14, x8: 12 ... x20: 3
This includes sites that are down.

dump(i) gives the committed values of all copies of all variables at site i in one line.

dump(xj) gives the committed values of all copies of variable xj at all sites, one line per site.

end(T1) causes your system to report whether T1 can commit in the format
T1 commits

or
T1 aborts

fail(6) says site 6 fails. (This is not issued by a transaction, but is just an event that the tester will execute.)

recover(7) says site 7 recovers. (Again, a tester-caused event) We discuss this further below.

A newline in the input means time advances by one. There will be one instruction per line.

Example (partial script with six steps in which transactions T1 commits, and one of T3 and T4 may commit)

begin(T1)
begin(T2)
begin(T3)
W(T1, x1,5)
W(T3, x2,32)
W(T2, x1,17) — will cause T2 to wait, but the write will go ahead after T1 commits
end(T1)
begin(T4)
W(T4, x4,35)
W(T3, x5,21)
W(T4,x2,21)
W(T3,x4,23)— T4 will abort because it's younger

Design

Your program should consist of two parts: a single transaction manager that translates read and write requests on variables to read and write requests on copies using the available copy algorithm described in the notes. The transaction manager never fails. (Having a single global transaction manager that never fails is a simplification of reality, but it is not too hard to get rid of that assumption by using a shared disk configuration.)

If the TM requests a read for transaction T and cannot get it due to failure, the TM should try another site (all in the same step). If no relevant site is available, then T must wait. This applies to read-only transactions as well which, for each variable x, must have access to the version of x that was the last to commit before the transaction begins. T may also have to wait for conflicting locks. Thus the TM may accumulate an input command

7

for T and will try it on the next tick (time moment). As mentioned above, while T is blocked (whether waiting for a lock to be released or a failure to be cleared), our test files will emit no new operations for T.

If a site fails and recovers, the DM would normally perform local recovery first (perhaps by asking the TM about transactions that the DM holds pre-committed but not yet committed), but this is unnecessary since, in the simulation model, commits are atomic with respect to failures (i.e., all writes of a committing transaction apply or none do). Therefore, all non-replicated variables are available for reads and writes. Regarding replicated variables, the site makes them available for writing, but not reading. In fact, a read for a replicated variable x will not be allowed at a recovered site until a committed write to x takes place on that site (see lecture notes on recovery when using the available copies algorithm).

During execution, your program should say which transactions commit and which abort and why. For debugging purposes you should implement (for your own sake) a command like querystate() which will give the state of each DM and the TM as well as the data distribution and data values. Finally, each read that occurs should show the value read.

4.1.1 Running the programming project

You will demonstrate the project to our very able graders. You will have 15-30 minutes to do so. The test should take a few minutes. The only times tests take longer are when the software is insufficiently portable. The version you send in should run on departmental servers or on your laptop.

The grader will ask you for a design document that explains the structure of the code. The grader will also take a copy of your source code for evaluation of its structure. Finally, the grader may ask you how you would change the code to achieve some new function.

4.1.2 Documentation and Code Structure

Because this class is meant to fulfill the large-scale programming project course requirement, please give some thought to design and structure. The design document submitted in late October should include all the major functions, their inputs, outputs and side effects. If you are working in a team, the author of the function should also be listed. That is, we should

see the major components, relationships between the components, and communication between components. This should work recursively within components. A figure showing interconnections and no more than three pages of text would be helpful.

The submitted code similarly should include as a header: the author (if you are working in a team of two), the date, the general description of the function, the inputs, the outputs, and the side effects.

In addition, you should provide a few testing scripts in plain text to the grader in the format above and say what you believe will happen in each test.

You should store you entire repository in git, enabling the graders to trace your progress. Please ignore this. We are not allowed to violate your privacy in 2018.

Finally, you will use reprozip and virtual machines to make your project reproducible even across different architectures for all time. Simply using portable Java or python or whatever is not the same. Packaging using reprozip and doing so correctly constitutes roughly 20% of the project grade.

4.2   Possibility 2 — Benchmarking/Tuning Project

Use everything you learn in class or from the tuning book and try to improve a real system that is available to you. Use substantial relations, e.g. 100 million rows and up. Specify the database management system, operating system, hardware platform including disks, memory size, and processor. For each attempted improvement, specify what happened to the system (whether it became faster or slower) and try to form a hypothesis as to why that occurred.

You should be prepared to discuss your project with me or potentially in front of the class.

4.3   Possibility 3 — Paper

This year, I invite papers on two topics:

1. A survey of the novel uses and challenges of transactional memory. (This might be helpful https://prezi.com/nbe9ryrueeeq/)

2. A survey of languages for time series databases.

If you want to write a paper though, it's a one person job and we need to discuss beforehand. Your model for a survey paper should be ACM Computing Surveys. Also, you should create four non-trivial problems based on the material you read and give the solutions.

5 Project Schedule — depends on project you choose

5.1 Schedule for Programming Project

Last class in September: Letter of intent that you are going to do programming project. Partner chosen if any. Please send this letter to
Anshul Sharma as10950@nyu.edu

Last class in October: design document. Please send this document to
Anshul Sharma as10950@nyu.edu

Project is due on December 4, 2018. Please send your code to the graders. Between December 5 and December 12, your project will be graded. You will make an appointment that week with the grader. We will figure out a randomized way to do this.

5.2 Schedule for Benchmarking/Tuning Project and for Scalaris/SQLite project

Last class in September: project outline (should fit on one page). Tuning problem you intend to address. System you plan to use and experimental question you plan to ask. This must be approved by me (Shasha) before you go on.

Last Class in October: status report. How are you doing? Any show-stoppers.

December 4, 2018: final report/demo to me.